

Tools and Algorithms for Deciding Timed Relations

Mihir Mehta

Department of Computer Science and Engineering,
Indian Institute of Technology, Delhi.
`cs1090197@cse.iitd.ac.in`

December 2012

Abstract

This is a report summarising the author's project on their B Tech Project for the academic year 2012-2013.

1 Objectives

- To develop a software toolkit that would enable users to verify various timed relations specifications and implementations expressed as timed automata.
 - To gain an understanding of the theory related to labeled transition systems, CCS processes and timed automata by surveying relevant literature.
 - To study tools already built by researchers for similar purposes.
 - To develop the software in a modular way with modules for language specification and modules for implementations of utility algorithms.
 - To implement algorithms for determining timed relations.

2 CCS processes

- A CCS process is an automaton with state and interfaces for interaction.
- The interaction is in the form of *actions* over communication ports known as *channels*.
- Given a port name a we refer to a as the label for input on the port and \bar{a} as the label for the output on the port.

- *Inaction*: This is the simplest CCS process, denoted by 0. No state transitions or communication can occur, in other words, this represents a deadlock.
- *Prefixing*: This is the simplest constructor; if P is a process and a is a label (input or output) then $a.P$ is also a process which can perform the action a in order to become the process P .
- *Naming*: We can give names to processes using syntax such as

$$N \stackrel{def}{=} a_1.a_2....0$$
This gives us the ability to define CCS processes recursively, such as this one:

$$\text{Parrot} \stackrel{def}{=} a.\bar{a}.\text{Parrot}$$
- *Choice*: If P and Q are processes, then $P + Q$ is a process as well which has the initial capabilities of both P and Q . The deadlock process 0 is the identity element for this, that is, $P + 0 = P$ is an identity.
- *Parallel Composition*: If P and Q are processes, then $P|Q$ is a process as well in which P and Q may proceed independently or communicate via complementary ports.
- *Restriction*: If P is a process and L is a set of channel names, then P/L is a process in which the component processes of P are the only processes which can communicate over channels from the set L .
- *Relabeling*: If P is a process and f is a function from labels to labels, then $P[f]$ is a process where each label from the domain of f is replaced by its image under f . One application of relabelling is the idea of *generic* processes: By relabelling the generic ports of such a process with specific port names, one can generate specific processes.

It is evident that each CCS process can be replaced by a labeled transition system (LTS) with equivalent behaviour, therefore we will, in the rest of this discussion, freely use the properties of LTS when describing those of CCS.

3 Equivalences on CCS

3.1 Trace equivalence

- A *trace* of an LTS is a sequence of actions that the LTS can perform.
- For an LTS P , the set $Traces(P)$ represents the set of all possible traces of P .
- Trace equivalence is said to exist between two LTS P and Q when $Traces(P) = Traces(Q)$.

- However, this notion proves to have a significant limitation in the case of CCS processes: Two CCS processes can have trace equivalence between their corresponding LTS and yet behave differently in terms of when they deadlock while interacting with a third CCS process.

3.2 Strong bisimilarity

- *Strong bisimulation*: A binary relation R is a strong bisimulation if and only if, for all $(s_1, s_2) \in R$ and $a \in Act$.

$$\forall s'_1 (s_1 \xrightarrow{a} s'_1 \Rightarrow \exists s'_2. (s_2 \xrightarrow{a} s'_2 \wedge (s'_1, s'_2) \in R)) \wedge$$

$$\forall s'_2 (s_2 \xrightarrow{a} s'_2 \Rightarrow \exists s'_1. (s_1 \xrightarrow{a} s'_1 \wedge (s'_1, s'_2) \in R))$$
- *Strong bisimilarity*: It can be shown that the union of all strong bisimulations over the set of states is a strong bisimulation. This binary relation is called strong bisimilarity, denoted by \sim .
- This mitigates one of the failings of trace equivalence as an equivalence relation: strong bisimilarity between two CCS processes ensures identical deadlock behaviour while interacting with a third CCS process.
- However, another limitation soon becomes apparent: if two CCS processes are to be strongly bisimilar, they must coincide even on the number and position of τ transitions in their traces. This is contrary to the semantics of CCS processes, as a τ transition is supposed to be private to a process and invisible to all other processes in its environment.

3.3 Weak bisimilarity

- *Weak bisimulation*: A binary relation R is a **weak bisimulation** if and only if, for all $(s_1, s_2) \in R$ and $a \in Act$.

$$\forall s'_1 (s_1 \xrightarrow{a} s'_1 \Rightarrow \exists s'_2. (s_2 \xRightarrow{a} s'_2 \wedge (s'_1, s'_2) \in R)) \wedge$$

$$\forall s'_2 (s_2 \xrightarrow{a} s'_2 \Rightarrow \exists s'_1. (s_1 \xRightarrow{a} s'_1 \wedge (s'_1, s'_2) \in R))$$
- It can be shown that the union of all weak bisimulations over the set of states is a weak bisimulation. This binary relation is called **weak bisimilarity**, denoted by \approx .
- Better suited to CCS processes, as it ignores τ transitions, thus disregarding hidden behaviour within a process.

3.4 Kanellakis and Smolka's algorithm

- This is a naive algorithm for determining the bisimilarity relation for the set of processes in a labelled transition system.
- This relies on the notion of a splitter.
- Let $\pi = \{B_0, \dots, B_k\}, k \geq 0$ be a partition of the set of states Pr in a labeled transition system.

- A splitter for a block $B_i \in \pi$ is a block $B_j \in \pi$ such that for some action $a \in Act$, some states in B_i have a -labelled transitions whose targets lie in B_j while other states in B_i do not.
- This suggests a refinement of π : replace block B_i with

$$B_i^1 = B_i \cap T_a^{-1}[B_j]$$

$$B_i^2 = B_i - B_i^1$$
- Refinements of this kind constitute the steps of this algorithm.
- The time complexity of this algorithm is $O(mn)$, since there can be at most n iterations, and all m edges are scanned in each iteration.

3.5 Fernandez' algorithm

- This is a more efficient algorithm for determining bisimilarity ($O(m \log n)$).
- This relies on the technique of three-way splitting introduced by Paige and Tarjan.
- Splitters can now be 'simple' or 'compound'.
- Stability: A partition π is said to be stable with respect to a compound block S if S is not a splitter for any block in π for any action.
- For a compound block S , having a constituent simple block B satisfying $n(B) \leq 0.5 * n(S)$, and with respect to which π is stable, we can split a block B_i on an action a as follows:

$$B_i^1 = (B_i \cap T_a^{-1}[B]) - T_a^{-1}[S - B]$$

$$B_i^2 = (B_i \cap T_a^{-1}[S - B]) - T_a^{-1}[B]$$

$$B_i^3 = B_i \cap T_a^{-1}[B] \cap T_a^{-1}[S - B]$$

4 Timed automata

- Formally, a timed automaton over a finite set of clocks C and a finite set of actions Act is a 4-tuple (L, l_0, E, I) .
- L is a finite set of locations.
- l_0 is the initial location.
- $E \subseteq L \times B(C) \times Act \times 2^C \times L$ is a finite set of edges.
- $I : L \rightarrow B(C)$ assigns invariants to each edge location.
- $B(C)$ is the set of clock constraints over C . An element of $B(C)$ can be an equality, a slack inequality, a strict inequality, or an AND combination of such constraints.

5 Equivalences on Timed Automata

5.1 Timed bisimilarity

- *Timed bisimulation*: A binary relation R is a timed bisimulation if and only if, for all $(s_1, s_2) \in R$, $a \in Act$, $d \in R_{\geq 0}$

$$\begin{aligned} \forall s'_1 (s_1 \xrightarrow{a} s'_1 \Rightarrow \exists s'_2. (s_2 \xrightarrow{a} s'_2 \wedge (s'_1, s'_2) \in R)) \wedge \\ \forall s'_2 (s_2 \xrightarrow{a} s'_2 \Rightarrow \exists s'_1. (s_1 \xrightarrow{a} s'_1 \wedge (s'_1, s'_2) \in R)) \wedge \\ \forall s'_1 (s_1 \xrightarrow{d} s'_1 \Rightarrow \exists s'_2. (s_2 \xrightarrow{d} s'_2 \wedge (s'_1, s'_2) \in R)) \wedge \\ \forall s'_2 (s_2 \xrightarrow{d} s'_2 \Rightarrow \exists s'_1. (s_1 \xrightarrow{d} s'_1 \wedge (s'_1, s'_2) \in R)) \end{aligned}$$
- It can be shown that the union of all timed bisimulations over the set of states is a timed bisimulation. This binary relation is called **timed bisimilarity**, denoted by \sim .

5.2 Time abstracted bisimilarity

- *Time abstracted bisimulation*: A binary relation R is a time abstracted bisimulation if and only if, for all $(s_1, s_2) \in R$, $a \in Act$, $d \in R_{\geq 0}$

$$\begin{aligned} \forall s'_1 (s_1 \xrightarrow{a} s'_1 \Rightarrow \exists s'_2. (s_2 \xrightarrow{a} s'_2 \wedge (s'_1, s'_2) \in R)) \wedge \\ \forall s'_2 (s_2 \xrightarrow{a} s'_2 \Rightarrow \exists s'_1. (s_1 \xrightarrow{a} s'_1 \wedge (s'_1, s'_2) \in R)) \wedge \\ \forall s'_1 (s_1 \xrightarrow{d} s'_1 \Rightarrow \exists (s'_2, d'). (s_2 \xrightarrow{d} s'_2 \wedge (s'_1, s'_2) \in R)) \wedge \\ \forall s'_2 (s_2 \xrightarrow{d} s'_2 \Rightarrow \exists (s'_1, d'). (s_1 \xrightarrow{d} s'_1 \wedge (s'_1, s'_2) \in R)) \end{aligned}$$
- It can be shown that the union of all time abstracted bisimulations over the set of states is a time abstracted bisimulation. This binary relation is called **time abstracted bisimilarity**, denoted by \sim_u .

5.3 Regions and region graphs

- *Equivalence of valuations*: Two valuations v and v' of a timed automaton are said to be equivalent ($v \equiv v'$) if and only if:
 - For each $x \in C$, either both $v(x)$ and $v'(x)$ are greater than c_x or $\lfloor v(x) \rfloor = \lfloor v'(x) \rfloor$
 - For each $x \in C$ such that $v(x) \leq c_x$, $frac(v(x)) = 0$ if and only if $frac(v'(x)) = 0$.
 - For all $x, y \in C$ such that $v(x) \leq c_x$ and $v(y) \leq c_y$, we have $frac(v(x)) \leq frac(v(y))$ if and only if $frac(v'(x)) \leq frac(v'(y))$.
- Under this notion of equivalence, an equivalence class is known as a *region*. The equivalence class containing a valuation v is denoted by $[v]_{\equiv}$.

- *Region graph*: With these definitions, we can define the region graph of a timed automaton as a LTS representation of the automaton where the action transitions are labelled with the same actions and the delay self-transitions are labelled with ε . Formally, the region graph of a timed automaton A with clock set C and action set Act is an LTS

$$T_r(A) = (Proc, Act \cup \{\varepsilon\}, \{\xrightarrow{a} \mid a \in Act \cup \{\varepsilon\}\})$$

where $Proc = \{(l, [v]_{\equiv}) \mid l \in L, v : C \rightarrow R_{\geq 0}\}$ (these states are called symbolic states)

The transitions are defined as follows:

- For each $a \in Act$, $(l, [v]_{\equiv}) \xrightarrow{a} (l', [v']_{\equiv})$ iff $(l, v) \xrightarrow{a} (l', v')$
- $(l, [v]_{\equiv}) \xrightarrow{\varepsilon} (l, [v']_{\equiv})$ if for some $d \in R_{\geq 0}$, $(l, v) \xrightarrow{d} (l, v')$

5.4 Zones and reachability graphs

- *Extended clock constraints*: these are described by the grammar
 $g ::= x \bowtie n \mid x - y \bowtie n \mid g_1 \wedge g_2$
The set of extended clock constraints is denoted by $B^+(C)$
- *Zones*: A zone is a set of valuations represented by an extended clock constraint.
 $Z = v \mid v \in g_Z$
- *Symbolic state*: For a timed automaton A , a symbolic state is an ordered pair of a location and a zone, that is, (l, Z) .
- *Reachability graph*: The reachability graph of a timed automaton is a directed graph in which the nodes are symbolic states and the edges are given by the *reachability relation*:
 - (l, Z') is reachable from (l, Z) if there exists a valuation belonging Z' which can be obtained from Z by applying a time delay which does not violate the invariant at l .
 - (l', Z') is reachable from (l, Z) if l has an action transition to l' and there exists a valuation in Z' which violates neither the guard of the transition nor the invariant at l' .